



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

Escola d'Enginyeria de Barcelona Est

Dimensionality Reduction with Principal Component Analysis (PCA)

Raúl Benítez

raul.benitez@upc.edu

Christian Mata

Christian.mata@upc.edu

DIMENSIONALITY REDUCTION

Big & High-Dimensional Data

- High-Dimensions = Lot of Features

Document classification

Features per document =
thousands of words/unigrams
millions of bigrams, contextual
information



Surveys -Netflix

480189 users x 17770 movies

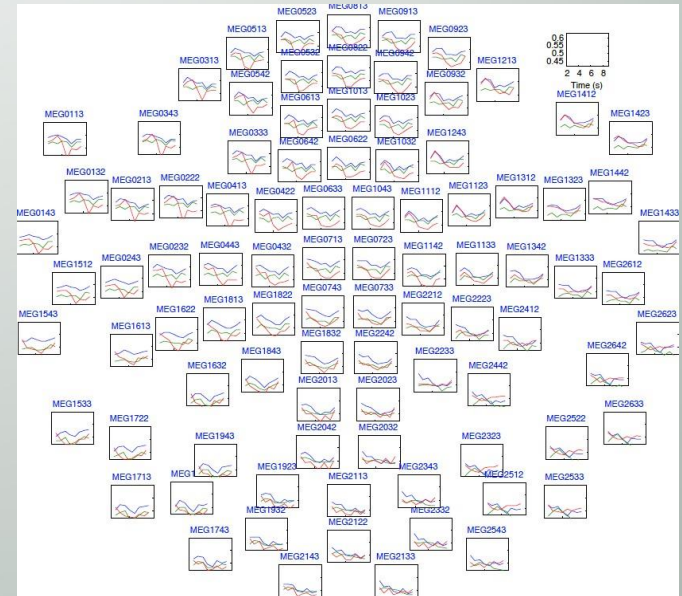

	movie 1	movie 2	movie 3	movie 4	movie 5	movie 6
Tom	5	?	?	1	3	?
George	?	?	3	1	2	5
Susan	4	3	1	?	5	1
Beth	4	3	?	2	4	2

Big & High-Dimensional Data

- High-Dimensions = Lot of Features

MEG Brain Imaging

120 locations x 500 time points
x 20 objects



Or any high-dimensional image data



Big & High-Dimensional Data

- Big & High-Dimensional Data.
- Useful to learn lower dimensional representations of the data.

Learning Representations

PCA, Kernel PCA, ICA: Powerful unsupervised learning techniques for extracting hidden (potentially lower dimensional) structure from high dimensional datasets.

Useful for:

- Visualization
- More efficient use of resources (e.g., time, memory, communication)
- Statistical: fewer dimensions → better generalization
- Noise removal (improving data quality)
- Further processing by machine learning algorithms


Principal Component Analysis (PCA)

Principal Component Analysis (PCA)

★ PART 1: The idea
What is PCA?

Principal Component Analysis (PCA)


What is PCA: Unsupervised technique for extracting variance structure from high dimensional datasets.




- PCA is an orthogonal projection or transformation of the data into a (possibly lower dimensional) subspace so that the variance of the projected data is maximized.

Principal Component Analysis (PCA)

Intrinsically lower dimensional than the dimension of the ambient space.




If we rotate data, again only one coordinate is more important.



Question:

Can we transform the features so that we only need to preserve one latent feature?

Principal Component Analysis (PCA)



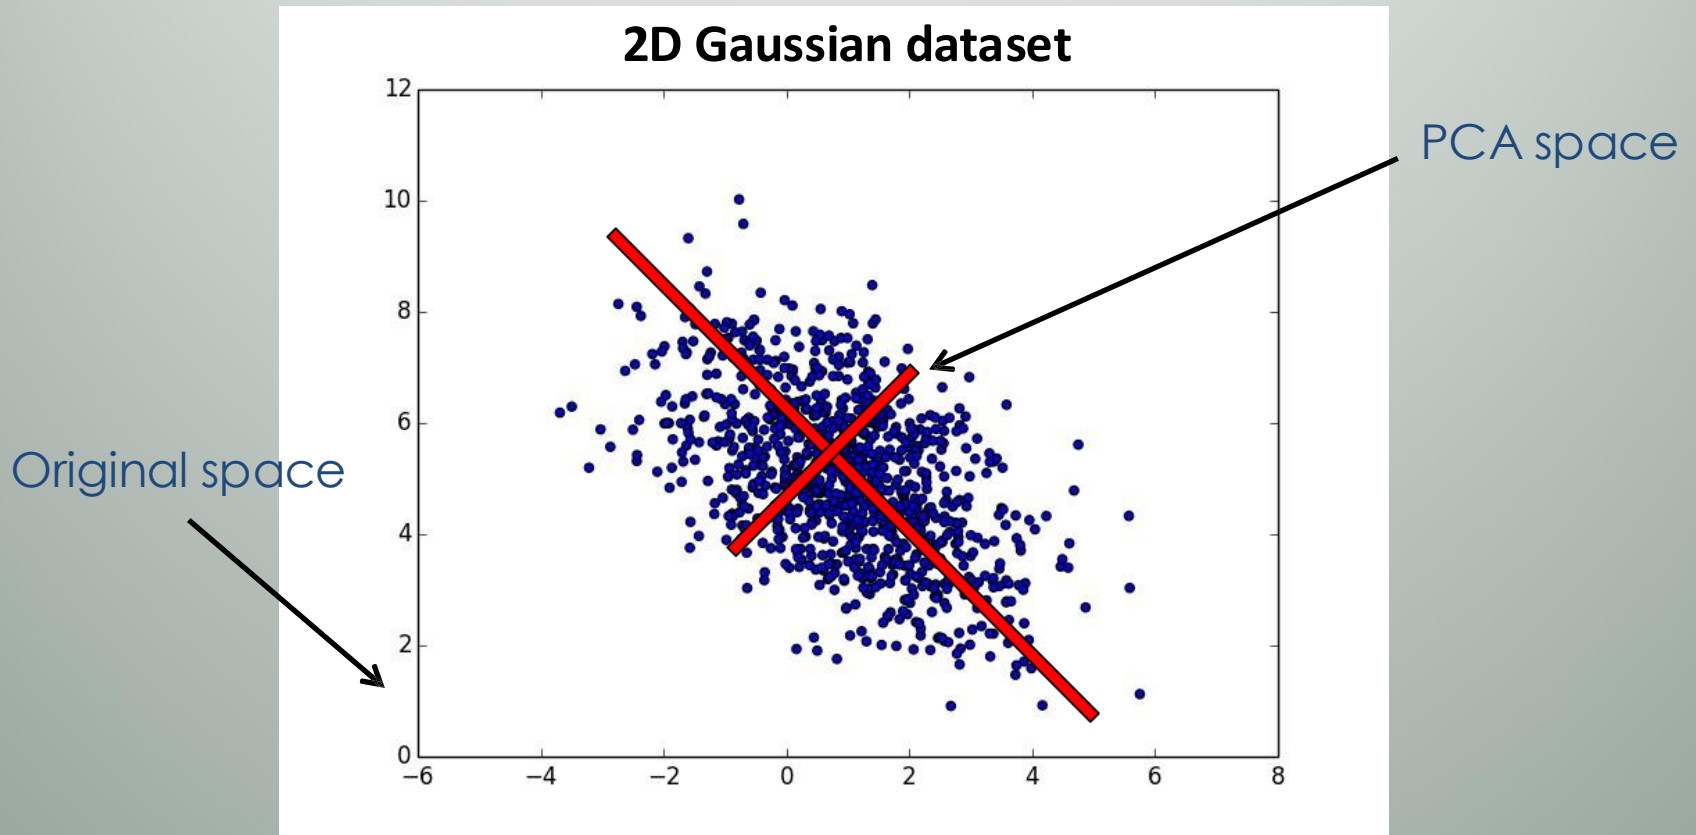
In case where data lies on or near a low d -dimensional linear subspace, axes of this subspace are an effective representation of the data.

Identifying the axes is known as [Principal Components Analysis](#), and can be obtained by using classic matrix computation tools (Eigen or Singular Value Decomposition).

Principal Component Analysis (PCA)

Dimensionality reduction:


- Simplify the description of your data
- Linear transformation of the original variables



Principal Component Analysis (PCA)

- Keep only the most relevant components (How many?)


2D Gaussian dataset



Principal Component Analysis (PCA)

- Keep only the most relevant components (How many?)
- Project each data point into the PCA subspace


2D Gaussian dataset



Principal Component Analysis (PCA)

- Keep only the most relevant components (How many?)
- Project each data point into the PCA subspace
- Use the projection as a new variable to describe your data


2D Gaussian dataset



Principal Component Analysis (PCA)

- Keep only the most relevant components (How many?)
- Project each data point into the PCA subspace
- Use the projection as a new variable to describe your data

2D Gaussian dataset



Principal Component Analysis (PCA)

✦ PART 1: The idea
What is PCA?

✦ PART 2: The math
How is it formulated?

Mathematical formulation:

Data: n variables, m observations

$$A_{m \times n} = \begin{matrix} & \text{variables} \\ \begin{pmatrix} x_1^1 & x_2^1 & \cdots & x_n^1 \\ x_1^2 & x_2^2 & \cdots & x_n^2 \\ x_1^3 & x_2^3 & \cdots & x_n^3 \\ \cdots & & & \\ x_1^m & x_2^m & \cdots & x_n^m \end{pmatrix} & \text{observations} \end{matrix}$$

Sample covariance matrix of data:

$$C_{n \times n} = (A - \bar{A})^T (A - \bar{A})$$

Procedure

1. Diagonalize covariance matrix:

$$C \cdot \vec{v}_i = \lambda_i \cdot \vec{v}_i, \quad i = 1 \dots n$$

Eigenvectors \vec{v}_i : Coordinates of the PCA space

Eigenvalues λ_i : Relevance of each PCA coordinate

Procedure

1. Diagonalize covariance matrix:

$$C \cdot \vec{v}_i = \lambda_i \cdot \vec{v}_i, \quad i = 1 \dots n$$

Eigenvectors \vec{v}_i : Coordinates of the PCA space

Eigenvalues λ_i : Relevance of each PCA coordinate

2. Keep PCA components by the variance of data explained by each component:

$$100 \leftarrow \frac{\sum_{i=1}^n \lambda_i}{\sum_{i=1}^n \lambda_i}$$

Procedure

1. Diagonalize covariance matrix:

$$C \cdot \vec{v}_i = \lambda_i \cdot \vec{v}_i, \quad i = 1 \dots n$$

Eigenvectors \vec{v}_i : Coordinates of the PCA space

Eigenvalues λ_i : Relevance of each PCA coordinate

2. Keep PCA components by the variance of data explained by each component:

$$100 \leftarrow \frac{\lambda_i}{\sum_{j=1}^n \lambda_j}$$

3. Project data into PCA subspace:

$$P_i = A \cdot \vec{v}_i$$

Projection of data into the PCA
coordinate associated to eigenvector \vec{v}_i

Example 1

Face recognition

Example 1: Face recognition

- Want to identify specific person, based on facial image
- Robust to glasses, lighting,...

Ⓜ Can't just use the given 256 x 256 pixels




Example 1: Face recognition

Method: Build one PCA database for the whole dataset and then classify based on the weights.



- Example data set: Images of faces
 - Famous Eigenface approach
[Turk & Pentland], [Sirovich & Kirby]

- Each face \mathbf{x} is ...
- 256 · 256 values (luminance at location)



– \mathbf{x} in 256·256 (view as 64K dim vector)

- Form $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_m]$ **centered** data mtx
- Compute $\mathbf{C} = \mathbf{X}\mathbf{X}^T$
- Problem: \mathbf{C} is 64K · 64K ... HUGE!!!

Example 1: Face recognition

Computational Complexity

- Suppose m instances, each of size N
 - Eigenfaces: $m=500$ faces, each of size $N=64K$
- Given $N \cdot N$ covariance matrix \mathbb{C} , can compute
 - all N eigenvectors/eigenvalues in $O(N^3)$
 - first k eigenvectors/eigenvalues in $O(k N^2)$
- But if $N=64K$, EXPENSIVE!

Example 1: Face recognition



Example 1: Face recognition

Reconstructing



- ... faster if train with...
 - only people w/out glasses
 - same lighting conditions

Example 1: Face recognition

Shortcomings

- Requires carefully controlled data:
 - All faces centered in frame
 - Same size
 - Some sensitivity to angle
- Alternative:
 - “Learn” one set of PCA vectors for each angle
 - Use the one with lowest error
- Method is completely knowledge free
 - (sometimes this is good!)
 - Doesn’t know that faces are wrapped around 3D objects (heads)
 - Makes no effort to preserve class distinctions

Example 2

Facial expression recognition

Example 2: Facial expression recognition

Happiness subspace




Example 2: Facial expression recognition

Disgust subspace



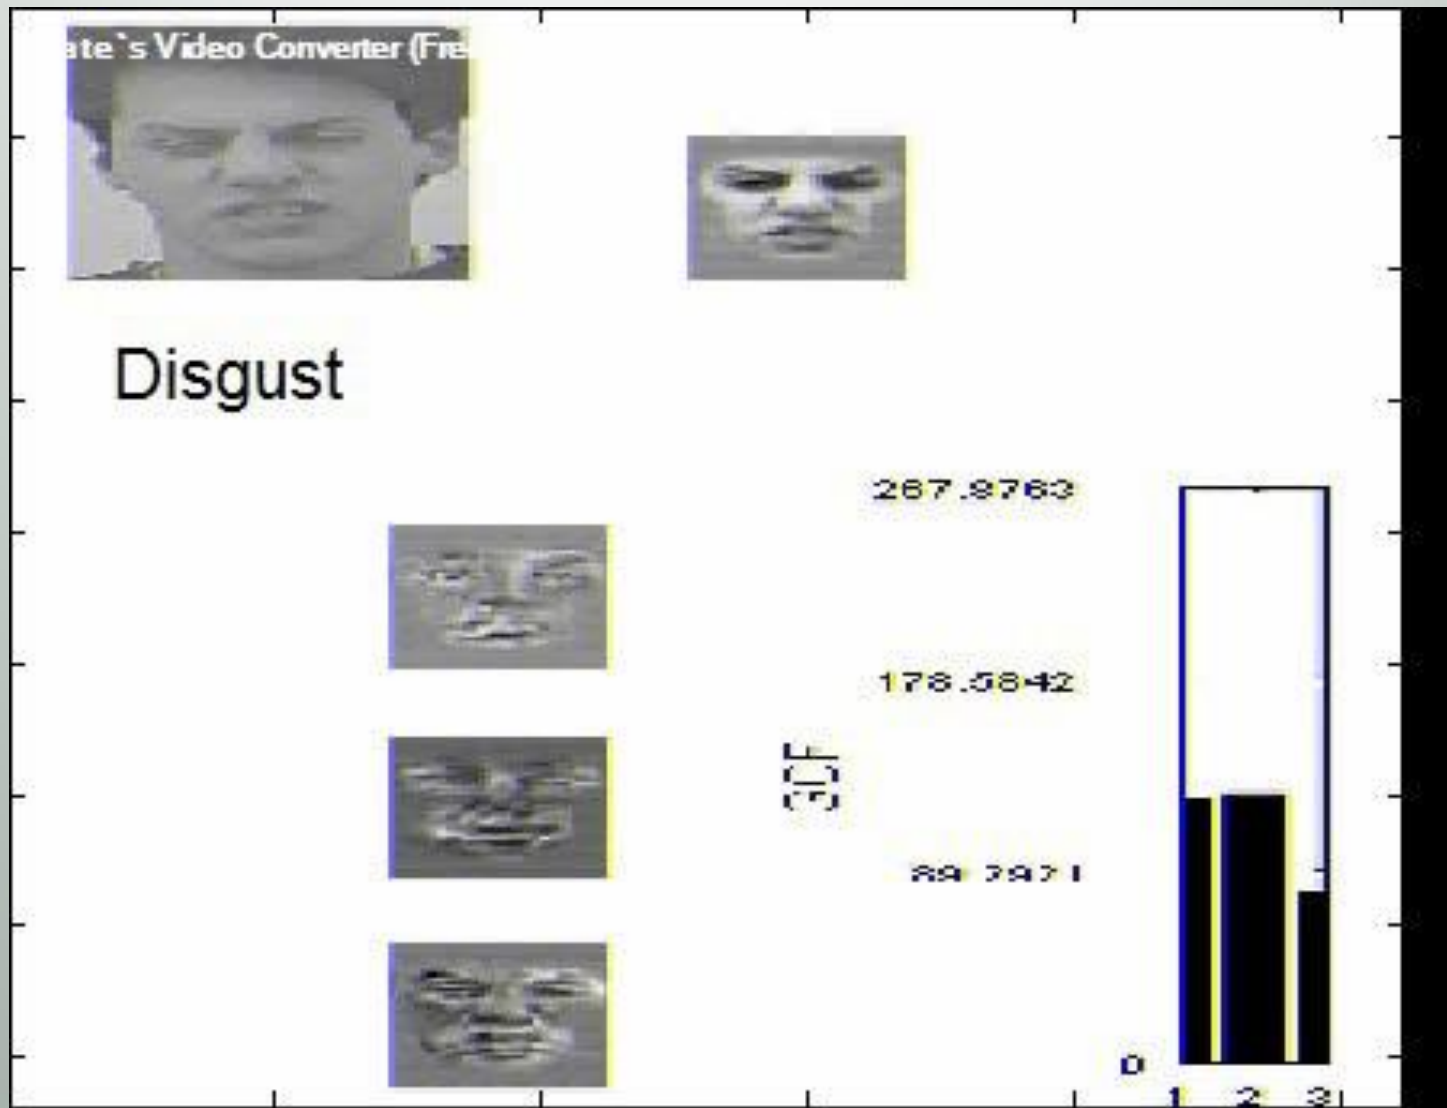
Example 2: Facial expression recognition



Example 2: Facial expression recognition



Example 2: Facial expression recognition



Example 3

Image compression

Example 3: Image compression


Original Image



- Divide the original 372x492 image into patches:
 - Each patch is an instance that contains 12x12 pixels on a grid
- View each as a 144-D vector

Example 3: Image compression

L_2 error and PCA dim



Example 3: Image compression

PCA compression: 144D \rightarrow 60D




Example 3: Image compression

PCA compression: 144D \rightarrow 16D



Example 3: Image compression

16 Most important eigenvectors





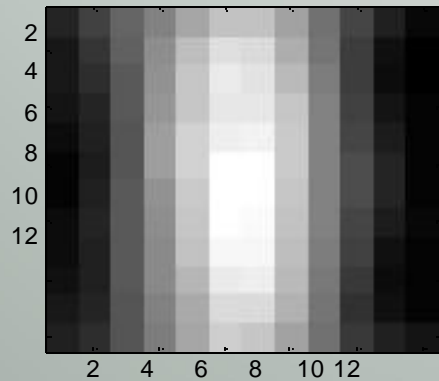



Example 3: Image compression

PCA compression: 144D \rightarrow 6D



Example 3: Image compression

6 Most important eigenvectors






Example 3: Image compression

PCA compression: 144D \rightarrow 3D




Example 3: Image compression

3 Most important eigenvectors




Example 3: Image compression

PCA compression: 144D \rightarrow 1D



Example 3: Image compression


60 Most important eigenvectors



Looks like the discrete cosine bases of JPG!...

Example 3: Image compression

2D Discrete Cosine Basis




Example 4

Noise filtering

Example 4: Noise filtering

Noise Filtering, Auto-Encoder...



Example 4: Noise filtering

Noisy image



Example 4: Noise filtering

Denoised image using 15 PCA components



Principal Component Analysis (PCA)

★ PART 1: The idea

What is PCA?

★ PART 2: The math

How is it formulated?

★ PART 3: The code

How to perform PCA in Python?

Python Implementation

Approach 1: do it manually (scipy)

```
21 # Obtain covariance matrix:
22 A1 = A - A.mean(0)
23 matcov = dot(A1.transpose(),A1)
24
25 # Diagonalization of covariance matrix:
26 valp,vecp = linalg.eig(matcov)
27
28 ind_creciente = argsort(valp) # sort eigenvalues
```

Approach 2: Use sklearn tools

```
46 # Use sklearn routines to obtain PCA projection:
47
48 from sklearn.decomposition import PCA
49
50 # Projection into 1d PCA space:
51 pca = PCA(n_components=1)
52 #X_r = pca.fit(A).transform(A)
53 X_r = pca.fit_transform(A)
```

More information

Numpy linear algebra-eigenvalues:

<http://docs.scipy.org/doc/numpy/reference/generated/numpy.linalg.eig.html>

Sklearn tools:

<http://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html>